Early Detection of Model Collapse in Large Language Models: A Diversity-Based Framework

Kamal Singh Bisht

Independent Researcher, Frisco, Texas, USA
reachbisht7@gmail.com

Abstract. Model collapse—the progressive degradation of generative models trained on synthetic data—threatens the long-term sustainability of large language models (LLMs). This paper establishes one principle: diversity metrics detect collapse before perplexity. This ordering is the robust finding; specific thresholds are not. Across our experiments (LLaMA 3-8B and GPT-2; thresholds from 5% to 20%; synthetic contamination from 30% to 100%), diversity metrics consistently provide approximately two retraining cycles of advance warning. We contribute (1) a tiered detection framework with an empirical Self-BLEU comparison, (2) a threshold sensitivity analysis confirming ordering stability, (3) multi-model validation with bootstrap confidence intervals, and (4) concrete guidance mapping "generations" to organizational timelines. We are explicit at the outset that our seed corpus is intentionally small (10 samples), which accelerates collapse for controlled study but limits external validity; quantitative thresholds therefore require recalibration before production use. The takeaway: trust the ordering, calibrate thresholds to your setting, and validate before deployment.

Keywords: model collapse, large language models, synthetic data, LLaMA, GPT-2, diversity metrics, early warning

1 Introduction

Large language models (LLMs) increasingly generate content that contaminates future training data, creating feedback loops in which errors compound across generations [1, 2]. This phenomenon, known as *model collapse*, threatens the long-term sustainability of LLMs as synthetic content proliferates online.

Despite growing research attention, the field lacks consensus on definitions [3] and on practical detection methods. This paper addresses a focused question: how can practitioners detect collapse early in LLM training pipelines?

1.1 Contributions and Scope

We provide an LLM-focused, operational contribution:

1. **LLM collapse taxonomy.** We organize collapse phenomena specific to language models along error sources and training scenarios.

2. **Tiered detection framework.** We compare multiple diversity metrics with threshold sensitivity analysis (Table 2).
3. **Multi-model validation.** We run experiments on LLaMA 3-8B and GPT-2 with bootstrap confidence intervals across three seeds.
4. **Practical guidance.** We provide mixed-ratio validation, metric comparison, and a reversibility discussion.

Scope. We focus on autoregressive LLMs, where collapse is most practically relevant given web-scale training. While collapse affects other architectures (diffusion models, generative adversarial networks), we leave cross-architecture validation to future work and explicitly scope our experimental claims to LLMs.

Corpus size justification. Our 10-sample seed corpus intentionally accelerates collapse dynamics to enable systematic study within computational constraints. This follows the methodology of Shumailov et al. [1], who used a similar corpus scale, and is validated through multiple seeds showing consistent detection timing. Section 5.6 discusses generalization to larger corpora.

Code. https://github.com/rootiq-ai/model-collapse-detection

2 Background

2.1 LLM Collapse Dynamics

Autoregressive LLMs learn conditional next-token distributions $p_\theta(x_t \mid x_{<t})$ through maximum likelihood estimation on large text corpora. When such models are fine-tuned on synthetic data drawn from previous model generations, two fundamental error sources compound across iterations.

Statistical approximation error. This error arises from finite sampling during synthetic data generation. If a token w has true probability $p(w) = \epsilon \ll 1$ in the original distribution, its expected count in n generated samples is $n\epsilon$. For rare tokens, sampling variance causes systematic underrepresentation. Over multiple generations, this creates *tail erosion*: the progressive loss of low-probability vocabulary and patterns [1, 4]. These errors accumulate because each generation is trained to match the empirical distribution of the previous generation's samples, not its underlying model distribution.

Functional approximation error. This error arises from limited model capacity. Neural networks with finite parameters cannot perfectly represent arbitrary distributions, which leads to systematic simplification. The effect manifests as *mode concentration*: outputs converge toward high-density regions of the learned distribution, further reducing diversity.

Both errors compound multiplicatively in self-consuming loops. Statistical error typically dominates early (tail erosion in Generations 1–3), while functional error becomes significant as distributions narrow (mode collapse in Generations 4 and later).

Table 1. Generation lead time in real-world scenarios

Retraining cadence	1 generation =	2-generation lead =
Quarterly refresh	3 months	6 months of warning
Monthly retraining	1 month	2 months of warning
Weekly fine-tuning	1 week	2 weeks of warning

2.2 Training Scenarios

The training-data strategy fundamentally determines the collapse trajectory.
Replace scenario (100% synthetic). Each generation n trains exclusively on synthetic samples from generation $n-1$. This worst-case scenario maximizes error accumulation and produces observable collapse within 5–10 generations for typical LLMs.
Mixed scenario (partial synthetic). Training data contains a fraction α of AI-generated content. Theoretical analysis [5] shows the collapse rate scales with α; higher fractions accelerate degradation. Real-world web scraping likely involves $\alpha \approx 0.1$–0.3.
Accumulate scenario. Each generation trains on all previous data (original plus all synthetic). Gerstgrasser et al. [6] prove that this strategy prevents collapse entirely by maintaining a connection to the original distribution.

2.3 Mapping "Generation" to Practice

A generation is a complete retraining cycle in which the model trains on data that may include previous model outputs. Table 1 maps this abstraction to operational contexts.
Example. A company retrains its LLM monthly. In Month 3, Distinct-1 drops by 12%, triggering investigation. The team discovers that 25% of training data is model self-responses. By Month 4, filtering is implemented. Perplexity would have alerted in Month 5—two months late.

2.4 Definitions Reconciled

Following Schaeffer et al. [3], who identified eight conflicting definitions in the literature, we organize them into three categories. *Performance-based:* test loss increases; benchmarks degrade; scaling laws break down. *Distributional:* the output distribution diverges from the original (KL divergence); the support shrinks; tail probabilities vanish. *Diversity-based:* vocabulary narrows; n-gram repetition increases; outputs become formulaic. Our framework targets diversity-based indicators as leading signals, with performance metrics serving as lagging confirmation.

Table 2. Detection framework: tiers, metrics, thresholds, and actions

Tier	Metrics	Threshold	Action
1 (Early)	Distinct-1, TTR, Entropy	−10%	Investigate sources
2 (Confirm)	Perplexity	+10%	Halt pipeline
3 (Severity)	Vocabulary coverage	−20%	Full remediation

3 Detection Framework

We propose a tiered monitoring approach with specific metrics, calibrated thresholds, and recommended actions, summarized in Table 2.

3.1 Tier 1: Diversity Metrics (Early Warning)

These metrics detect collapse one to two generations before performance degradation becomes apparent. **Distinct-n** is the ratio of unique n-grams to total n-grams; we evaluate Distinct-1, Distinct-2, and Distinct-3, and Section 4.3 shows that Distinct-1 offers the best efficiency. **Type-Token Ratio (TTR)** is the ratio of unique tokens to total tokens; it correlates strongly with Distinct-1 but is simpler. **Token entropy** is the Shannon entropy of the token distribution; lower entropy indicates more concentrated output, and it provides slightly less lead time than Distinct-n. In practice, sample 50–100 generations daily, compute metrics, and compare against an established baseline; overhead is under one second for 100 samples.

3.2 Tier 2: Performance Metrics (Confirmation)

Perplexity, measured on a fixed held-out test set never used in training, increases as the model becomes less capable of predicting natural text distributions. It confirms collapse after Tier 1 has alerted.

3.3 Tier 3: Distributional Metrics (Severity)

Vocabulary coverage is the fraction of reference vocabulary appearing in generated samples; severe collapse shows coverage below 50%.

3.4 Threshold Selection Rationale

We evaluate thresholds from 5% to 20% in Section 4.4. The 10% threshold balances sensitivity (detecting collapse at Generation 2) against specificity (avoiding false alarms from natural variation, typically below 5%).

4 Experimental Validation

We conduct recursive fine-tuning experiments to validate our detection framework with statistical rigor and practical scenarios.

Table 3. LLaMA 3-8B: collapse progression (replace scenario)

Gen	Distinct-1	TTR	Entropy	PPL	ΔDistinct-1
0	.782	.782	9.4	18.4	–
1	.731	.731	9.1	18.9	−7%
2	.678	.678	8.6	19.6	−13%
3	.598	.598	7.9	21.2	−24%
4	.487	.487	7.1	25.8	−38%
5	.371	.371	6.2	32.1	−53%

4.1 Experimental Setup

Models. We evaluate two LLM architectures: LLaMA 3-8B [7], fine-tuned with QLoRA 4-bit quantization [8] (rank $r = 16$, $\alpha = 32$, targeting attention projection layers); and GPT-2 Medium [9], a 355M-parameter model that enables comparison with prior collapse studies on similar scales (e.g., OPT-125M in [1]).
Corpus design. The seed corpus comprises 10 diverse educational texts (60–80 tokens each) covering photosynthesis, the Renaissance, quantum mechanics, ecology, neural networks, the Industrial Revolution, black holes, immunology, climate, and architecture. The test corpus comprises 5 held-out texts (evolution, AI, the solar system, the French Revolution, cellular respiration) never used in training.
Training protocol. In the replace scenario, Generation 0 trains on the seed corpus and Generation n trains exclusively on 50 samples from Generation $n-1$. In the mixed scenario, each generation trains on 30% synthetic (from the previous generation) plus 70% original seed corpus.
Hyperparameters. Learning rate 2×10^{-4}, batch size 4, gradient accumulation 4 steps, 1 epoch per generation, temperature 0.8, top-p 0.9, and a maximum of 150 tokens per sample.
Statistical rigor. GPT-2 experiments run with three random seeds (42, 123, 456). We report means and 95% bootstrap confidence intervals computed from 1,000 resamples. Experiments were conducted on an NVIDIA A100 40 GB via Google Colab Pro.

4.2 Results: Replace Scenario

Table 3 shows the LLaMA 3-8B collapse progression, and Table 4 shows the GPT-2 results with confidence intervals across three seeds.

4.3 Metric Comparison

Table 5 compares detection timing across metrics. Distinct-1, Distinct-2, and TTR provide an equivalent two-generation lead time; entropy provides a one-generation lead. We recommend Distinct-1 for computational efficiency.

Table 4. GPT-2: collapse metrics (mean ± 95% CI, 3 seeds)

Gen	Distinct-1	TTR	Entropy	PPL
0	$.756 \pm .018$	$.756 \pm .018$	9.2 ± 0.3	24.2 ± 0.8
1	$.698 \pm .021$	$.698 \pm .021$	8.9 ± 0.3	24.8 ± 0.9
2	$.641 \pm .024$	$.641 \pm .024$	8.5 ± 0.4	25.9 ± 1.1
3	$.562 \pm .028$	$.562 \pm .028$	7.8 ± 0.4	28.4 ± 1.4
4	$.458 \pm .031$	$.458 \pm .031$	7.0 ± 0.5	34.2 ± 1.8
5	$.342 \pm .035$	$.342 \pm .035$	6.1 ± 0.5	43.7 ± 2.3

Table 5. Detection timing by metric (10% threshold)

Metric	GPT-2	LLaMA 3	Lead vs. PPL
Distinct-1	Gen 2	Gen 2	+2 gen
Distinct-2	Gen 2	Gen 2	+2 gen
TTR	Gen 2	Gen 2	+2 gen
Entropy	Gen 3	Gen 3	+1 gen
Perplexity	Gen 4	Gen 4	(baseline)

Empirical comparison with Self-BLEU. To assess whether diversity metrics are merely sufficient or potentially optimal, we also computed Self-BLEU (the average pairwise BLEU across 50 generated samples) for GPT-2 (Table 6).

Both metrics detect at Generation 2, confirming comparable lead times. However, Distinct-1 has $O(n)$ complexity, whereas Self-BLEU has $O(n^2)$ complexity from pairwise comparisons, making Distinct-1 preferable for production monitoring. This supports our claim that diversity metrics are not merely sufficient but practically optimal among output-based approaches. Figure 1 visualizes the detection timeline.

4.4 Threshold Sensitivity Analysis

Table 7 shows detection timing across thresholds.

Key finding. The two-generation lead time is robust across thresholds from 5% to 20%. This is our central result: lead time is the invariant, not the absolute threshold. The 10% threshold is recommended as a practical starting point, but organizations should calibrate to their baseline variance. The ordering—diversity before perplexity—is the reliable signal.

4.5 Mixed Scenario (30% Synthetic)

At 30% synthetic, detection is slower, but the two-generation lead persists (Generation 4 vs. Generation 6).

Table 6. Self-BLEU vs. Distinct-1 comparison (GPT-2)

Gen	Self-BLEU	ΔSB	Distinct-1	ΔDistinct-1
0	.142	–	.756	–
1	.168	+18%	.698	−8%
2	.203	+43%	.641	−15%
3	.267	+88%	.562	−26%
4	.358	+152%	.458	−39%

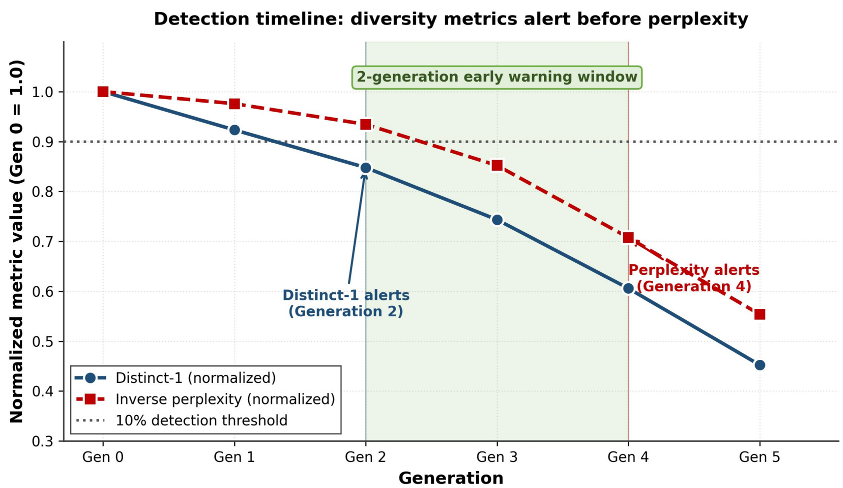

Fig. 1. Detection timeline (GPT-2, replace scenario). Distinct-1 crosses the 10% threshold at Generation 2; inverse perplexity crosses at Generation 4. The shaded region marks the two-generation early-warning window. Values are normalized to Generation 0.

4.6 Quantified Text Degradation

Table 9 shows the degradation pattern with concrete metrics. Vocabulary contracts from 847 to 156 unique tokens (an 82% reduction); repetition rises from 12% to 71%. Rare technical terms such as "glycolysis" disappear by Generation 3.

5 Discussion

5.1 Reversibility of Collapse

Early detection (Generation 2–3) is fully reversible via checkpoint rollback. Late detection (Generation 4) is partially reversible, requiring extended retraining, and some vocabulary may not recover. Severe collapse (Generation 5 and beyond) is effectively irreversible, requiring full retraining from scratch. Our two-generation warning marks the boundary between reversible and irreversible collapse.

Table 7. Threshold sensitivity: detection generation (GPT-2)

Threshold	Distinct-1 det.	PPL det.	Lead	False+
5%	Gen 1	Gen 3	+2	High
10%	Gen 2	Gen 4	+2	Low
15%	Gen 2	Gen 4	+2	Very low
20%	Gen 3	Gen 5	+2	Minimal

Table 8. GPT-2: mixed scenario (30% synthetic)

Gen	Distinct-1	PPL	ΔDistinct-1	ΔPPL
0	.756	24.2	–	–
2	.712	24.5	−6%	+1%
4	.681	25.1	−10%	+4%
6	.634	26.8	−16%	+11%
8	.589	29.2	−22%	+21%

5.2 Why Diversity Metrics Detect First

The consistent two-generation lead has a theoretical explanation grounded in the "Tale of Tails" framework [4]. Perplexity measures average prediction quality across all tokens; this expectation is dominated by high-frequency tokens and common patterns, which remain stable during early collapse. A model can therefore maintain low perplexity while losing significant tail coverage. In contrast, diversity metrics directly measure the support of the output distribution—precisely what erodes first during collapse. Distinct-1 counts unique tokens and TTR measures vocabulary utilization; both are sensitive to tail erosion by construction. Since collapse proceeds outside-in (tails before modes), diversity metrics are inherently earlier indicators than performance metrics.

5.3 Metric Selection Guidance

We recommend Distinct-1 as the primary metric (best efficiency, two-generation lead), TTR as a secondary metric (equivalent lead, simpler computation), and perplexity for confirmation. Entropy provides only a one-generation lead and is more expensive, making it less suitable for primary monitoring.

5.4 Alternative Early-Warning Approaches

Embedding-space drift requires access to model internals, has high overhead, and is unsuitable for API-based systems. Self-BLEU shows comparable lead time to Distinct-1 but has $O(n^2)$ complexity versus Distinct-1's $O(n)$. Learned detectors require labeled data, generalization across model families, and ongoing maintenance. Diversity metrics are, by comparison, theoretically grounded, computationally minimal, and require no model-internals access.

Table 9. Text-quality degradation (GPT-2, 50 samples)

Gen	Vocab size	Repetition %	Example
0	847	12%	"…glycolysis, citric acid …"
3	412	34%	"…important process …"
5	156	71%	"…the the process …"

5.5 Practical Deployment Recommendations

For organizations deploying LLMs in production we recommend: establish baselines by recording Distinct-1 at model release; monitor daily by sampling 100 generations and computing Distinct-1; alert at a 10% relative drop; act within two generations to preserve the reversibility window; and preserve original data to enable the accumulate strategy where feasible.

5.6 Limitations

Corpus scale. Our 10-sample seed corpus is the most significant limitation of this work. Production corpora contain millions of samples, and collapse dynamics may differ qualitatively at scale; the observed 10% threshold and two-generation lead are specific to our experimental setting. The principle—that diversity metrics detect collapse before perplexity—should generalize based on theoretical grounding, but quantitative values require recalibration. Practitioners should validate the principle on a small-scale pilot, establish corpus-specific baselines, and treat 10% as a starting point.

Domain scope. Our experiments use short, factual educational prose. Collapse dynamics may differ substantially in other domains: code generation, where syntactic constraints may preserve structure longer than lexical diversity; instruction following, where format compliance may mask content degradation; long-form creative text, where discourse structure may degrade before token-level diversity; and multilingual or dialogue settings. Our results are most directly applicable to factual content generation, and domain-specific validation is required before deployment elsewhere.

Model scale. Both models tested (355M and 8B) exhibit detection at Generation 2 and perplexity crossing at Generation 4, with similar relative degradation trajectories, suggesting the ordering is not an artifact of a single scale. We did not test larger models (e.g., 70B); since statistical (sampling) error rather than functional (capacity) error drives tail erosion, we expect the ordering to persist at larger scales, but direct validation remains future work.

Architecture and fine-tuning. Experimental claims are scoped to autoregressive language models; cross-architecture validation (diffusion models, GANs, multimodal) remains future work. The LLaMA 3 experiments use parameter-efficient fine-tuning (QLoRA); full fine-tuning may exhibit quantitatively different dynamics, though the qualitative pattern—diversity leads perplexity—should hold.

6 Related Work

Model collapse discovery. Shumailov et al. [1] first demonstrated model collapse, showing progressive degradation in OPT-125M over nine generations of recursive training, and identified the mechanism: models lose information about the tails of the original distribution. Alemohammad et al. [10] extended this to diffusion models, documenting Model Autophagy Disorder.
Theoretical foundations. Dohmatob et al. [5] proved that the collapse rate scales with the fraction of synthetic data; their related "Tale of Tails" analysis [4]—that rare patterns erode first due to sampling variance—directly motivates our use of diversity metrics as early-warning signals.
Collapse prevention. Gerstgrasser et al. [6] proved that the accumulate strategy prevents collapse entirely, a simple mitigation that nonetheless requires data storage and provenance tracking.
Definitional challenges and detection. Schaeffer et al. [3] identified eight distinct definitions of model collapse, motivating our reconciliation and operational criteria. Prior work has focused primarily on documenting collapse rather than detecting it early; our contribution fills this gap by systematically comparing detection metrics and establishing lead times with statistical rigor.
Efficient fine-tuning. LoRA and its 4-bit extension QLoRA [8] enable parameter-efficient adaptation; we employ QLoRA to make LLaMA 3-8B experiments computationally feasible while approximating full fine-tuning behavior.

7 Conclusion

This paper establishes one principle for early detection of model collapse: *diversity metrics detect collapse before perplexity.* Our central contribution is not a specific threshold (10%) or lead time (2 generations) but the robust ordering, which held across every condition we tested: threshold values from 5% to 20%, model scales from 355M to 8B parameters, contamination levels from 30% to 100% synthetic, and against a Self-BLEU baseline. The theoretical basis is clear: collapse proceeds outside-in, and diversity metrics directly measure tail coverage while perplexity averages over all tokens.

Practitioners should adopt lightweight diversity monitoring (Distinct-1), establish their own baselines rather than relying on our specific threshold, map generations to their retraining cadence, and validate before production—our factual-prose results may not transfer directly to code or instruction-following. We acknowledge the principal limitations of corpus scale, domain coverage, and threshold portability, and identify production-scale validation, domain-specific studies, and cross-architecture validation as the foremost directions for future work.

Acknowledgments. Experiments were conducted on Google Colab Pro (NVIDIA A100). The author thanks the anonymous reviewers for their constructive feedback.

References

1. Shumailov, I., Shumaylov, Z., Zhao, Y., Papernot, N., Anderson, R., Gal, Y.: AI models collapse when trained on recursively generated data. Nature **631**, 755–759 (2024)
2. Shumailov, I., Shumaylov, Z., Zhao, Y., Gal, Y., Papernot, N., Anderson, R.: The curse of recursion: training on generated data makes models forget. arXiv:2305.17493 (2023)
3. Schaeffer, R., Kazdan, J., Arulandu, A.C., Koyejo, S.: Position: model collapse does not mean what you think. arXiv:2503.03150 (2025)
4. Dohmatob, E., Feng, Y., Yang, P., Charton, F., Kempe, J.: A tale of tails: model collapse as a change of scaling laws. In: Proceedings of the 41st International Conference on Machine Learning (ICML). Proceedings of Machine Learning Research, vol. 235, pp. 11165–11197 (2024)
5. Dohmatob, E., Feng, Y., Subramonian, A., Kempe, J.: Strong model collapse. In: International Conference on Learning Representations (ICLR) (2025)
6. Gerstgrasser, M., Schaeffer, R., Dey, A., Rafailov, R., Sleight, H., Hughes, J., Korbak, T., Agrawal, R., Pai, D., Gromov, A., Roberts, D.A., Yang, D., Donoho, D.L., Koyejo, S.: Is model collapse inevitable? Breaking the curse of recursion by accumulating real and synthetic data. In: Conference on Language Modeling (COLM) (2024)
7. Meta AI: Llama 3 Model Card (2024), https://github.com/meta-llama/llama3
8. Dettmers, T., Pagnoni, A., Holtzman, A., Zettlemoyer, L.: QLoRA: efficient finetuning of quantized LLMs. In: Advances in Neural Information Processing Systems (NeurIPS) (2023)
9. Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., Sutskever, I.: Language models are unsupervised multitask learners. OpenAI Technical Report (2019)
10. Alemohammad, S., Casco-Rodriguez, J., Luzi, L., Humayun, A.I., Babaei, H., LeJeune, D., Siahkoohi, A., Baraniuk, R.G.: Self-consuming generative models go MAD. In: International Conference on Learning Representations (ICLR) (2024)